

Matrix-Residual Graph Neural Networks for Graph Classification

Xuelin Hu¹ and Xiaoqin Fu²

¹School of Railway Communication and Signaling Engineering, Liuzhou Railway Vocational Technical College, Liuzhou 545000, China

²School of Railway Locomotive and Rolling Stock, Liuzhou Railway Vocational Technical College, Liuzhou 545000, China

Corresponding author:
Xiaoqin Fu²

Email address: fuxiaoqin@ltzy.edu.cn

ABSTRACT

Background. Residual connections are widely used to stabilize graph neural networks, but residual design is often treated as a one-dimensional depth-wise choice. Multi-branch graph classifiers create a richer design space: residual information can flow along depth, across branches, or through a local two-dimensional branch-by-layer neighborhood.

Methods. We study Matrix-Residual Graph Neural Networks as a controlled residual-topology framework for graph classification. The model family represents hidden states on a branch-by-layer grid and compares five information-flow designs: plain message passing, vertical residual reuse, horizontal branch-wise exchange, local matrix residual reuse, and sparse/gated matrix reuse. The main benchmark evaluates six datasets, four message-passing operators, five model families, and five graph-level folds under a shared training protocol. We further run branch-count ablations, fold-0 sensitivity scans, five-fold checks of selected MatrixResGated candidates, and mechanism analyses. To further test effectiveness across structural granularities, we construct a controlled synthetic multi-class benchmark spanning five graph families and two to eight structural classes.

Results. Real-data experiments show that no single residual topology is universally optimal across datasets and operators. Branch-count ablations reveal dataset-dependent operating regions rather than monotonic capacity gains. The synthetic multi-class benchmark further shows that simple baselines remain competitive on low-class tasks, whereas MatrixRes and MatrixResGated become more consistently favorable as structural class granularity increases.

Conclusions. Residual topology, branch count, and residual filtering should be selected jointly. Additional branches help only when they create useful functional diversity; otherwise, residual traffic and branch redundancy can offset the added capacity. Matrix-style reuse is best understood as a tunable residual family rather than a universal replacement for simpler vertical or horizontal residual paths, with its clearest advantage appearing under finer-grained structural discrimination.

1 INTRODUCTION

Graph neural networks (GNNs) learn graph representations by repeatedly aggregating information from local neighborhoods (Gori et al., 2005; Kipf and Welling, 2017). This message-passing view has become a standard tool for graph classification because it combines node attributes, local connectivity, and graph-level structure in a single differentiable model (Gilmer et al., 2017; Wu et al., 2021). In molecular and protein-oriented benchmarks, graph labels may depend on both local motifs and global organization, so useful intermediate evidence must survive until graph-level pooling (Duvenaud et al., 2015; Ying et al., 2018).

Increasing depth or adding parallel branches can expand model capacity, but it also changes information transport through the network. Deep GNNs may suffer from over-smoothing, where node representations become too similar across propagation steps (Li et al., 2018; Oono and Suzuki, 2020). They may also suffer from over-squashing, where long-range information is compressed into limited-size messages (Alon and Yahav, 2021; Topping et al., 2021). These

issues are especially relevant for graph-level prediction because the final readout depends on the quality of the intermediate representations that reach the pooling stage.

Residual connections are a common response to these difficulties. They improve gradient flow and allow earlier representations to be reused, a principle established in deep convolutional networks and later adapted to graph architectures (He et al., 2016; Bresson and Laurent, 2017). Other depth-control strategies, including DropEdge, Jumping Knowledge aggregation, and graph-specific normalization, also aim to preserve useful information across propagation steps (Rong et al., 2019; Xu et al., 2018). However, residual connectivity is still commonly treated as a default layer-wise component rather than as the object of study. Most residual graph classifiers reuse information along depth in the same stream.

Once a model contains multiple branches, depth-wise residual reuse is only one possible route. Residual information may also move across neighboring branches at the same depth or through diagonal branch-depth neighborhoods. Existing graph-classification work has explored stronger message-passing operators, hierarchical pooling, attention, and higher-order neighborhood mixing (Abu-El-Haija et al., 2019; Bianchi et al., 2020). These designs improve the expressive or aggregation capacity of GNNs, but they do not isolate the residual-routing question: if several branch states are available at a given depth, should the model reuse only the previous state of the same branch, exchange states laterally, or combine both directions in a local two-dimensional residual neighborhood?

This paper studies residual reuse as a structured design problem on a branch-by-layer grid. The central question is not whether skip connections should be present, but where residual signals should flow. A grid view exposes three controlled neighborhoods: vertical reuse along depth, horizontal exchange across neighboring branches, and matrix-style reuse that combines same-branch, neighboring-branch, and diagonal residual paths. A sparse/gated matrix variant further tests whether matrix residual traffic should be filtered before injection. This framing follows controlled benchmark practice: the goal is not to add many unrelated architectural components, but to make one design axis explicit enough to evaluate (Keriven et al., 2020; Morris et al., 2020).

The study is deliberately controlled. It fixes the graph-classification pipeline within each comparison, evaluates the same architecture family under several canonical message-passing operators, and changes only the residual topology and residual-control mechanism. This separation keeps the empirical comparison focused on information-flow design rather than on unrelated changes in readout structure, task-specific feature engineering, or operator-specific tuning. Recent graph models may achieve stronger absolute performance by changing several components simultaneously; here, those components are intentionally held fixed so that the residual topology itself remains identifiable.

Real graph benchmarks are also limited in the range of structural class granularity that they expose. A residual topology may look adequate on binary graph classification while behaving differently when the label space requires finer structural separation. To test this point without adding task-specific feature engineering, we complement the real-data benchmark with controlled synthetic graph families whose class labels are generated by structural parameters. This synthetic setting is not intended to replace real benchmarks; it is a stress test for whether residual routing remains stable as structural classes become more fine-grained.

The work is organized around three research questions:

1. Does residual topology affect performance under a shared graph-classification protocol? We compare plain message passing, vertical residual reuse, horizontal branch-wise exchange, local matrix residual reuse, and sparse/gated matrix reuse while holding the remaining pipeline fixed.
2. How do branch count, residual traffic, and mechanisms explain performance changes? We test whether additional branches improve accuracy monotonically or only within a useful operating region, and we relate these trends to diversity, similarity, gradients, and residual traffic.
3. When is matrix-style residual reuse more effective on real data and controlled synthetic multi-class structures? We test whether matrix residual traffic should be passed densely or

102 filtered, whether candidate residual-control settings remain valid after full-fold reruns, and
 103 whether matrix-style reuse becomes more stable as structural class granularity increases.

104 The main contributions are:

- 105 • A branch-by-layer formulation of residual graph classifiers. Vertical, horizontal, matrix,
 106 and sparse/gated matrix residual families are expressed as local neighborhoods on the
 107 same grid, making their information-flow differences explicit.
- 108 • A controlled real-data benchmark and synthetic multi-class benchmark. The experiments
 109 keep the classifier, readout, fold protocol, and backbone choices matched, and add a struc-
 110 tural stress test spanning five controlled graph families and two to eight classes.
- 111 • A mechanism analysis with conservative rerun checks. Accuracy is analyzed together
 112 with branch diversity, branch cosine similarity, branch CKA, gradient norms, and residual
 113 activity. Sensitivity scans are used only to identify candidates, and selected settings are
 114 rerun across folds before being treated as empirical evidence.

115 2 MATERIALS AND METHODS

116 This section defines the graph-classification task, the Matrix-Residual architecture family, the
 117 benchmark protocol, the synthetic structural stress test, and the mechanism metrics. The
 118 guiding principle is to isolate residual topology while holding the remaining learning pipeline
 119 fixed.

120 2.1 Problem definition and notation

121 Let $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ be a graph with node feature matrix $\mathbf{X} \in \mathbb{R}^{n \times d_{\text{in}}}$ and adjacency structure $\mathbf{A}$. For
 122 a dataset $\mathcal{D} = \{(\mathcal{G}_i, y_i)\}_{i=1}^N$, the objective is to learn $f_{\theta} : \mathcal{G} \rightarrow \hat{y}$, where $\hat{y}$ is the predicted graph
 123 label.

124 The compared model families share the same graph-classification wrapper:

- 125 • Plain: stacked message passing without explicit residual reuse.
- 126 • VerticalRes: same-branch residual reuse from the previous layer.
- 127 • HorizontalRes: cross-branch exchange between adjacent branches at the same depth.
- 128 • MatrixRes: local two-dimensional reuse over vertical, horizontal, and diagonal neighbors.
- 129 • MatrixResGated: matrix reuse with gated or sparse residual filtering.

130 2.2 Shared graph-classification pipeline

131 Let B be the number of branches and L be the number of message-passing layers. The hidden
 132 state at branch b and layer ℓ is written as $\mathbf{H}_b^{(\ell)}$, where $b \in \{1, \dots, B\}$. A branch-local propagation
 133 step is

$$\mathbf{Z}_b^{(\ell)} = \phi_b^{(\ell)} \left(\mathbf{H}_b^{(\ell-1)}, \mathbf{A} \right), \quad (1)$$

134 where ϕ is instantiated as one of four standard operators: GCNConv (Kipf and Welling, 2017),
 135 GATConv (Veličković et al., 2018), SAGEConv (Hamilton et al., 2017), or GINConv (Xu et al.,
 136 2019). The residual-enhanced update is

$$\mathbf{H}_b^{(\ell)} = \mathbf{Z}_b^{(\ell)} + \sum_{(b', \ell') \in \mathcal{N}(b, \ell)} \mathcal{R}_{(b', \ell') \rightarrow (b, \ell)} \left(\mathbf{H}_{b'}^{(\ell')} \right). \quad (2)$$

137 The only term that changes across architectures is the residual neighborhood $\mathcal{N}(b, \ell)$ and the
 138 residual transform $\mathcal{R}$.

139 After the final message-passing layer, each branch is pooled to a graph-level vector by global
 140 mean pooling and the branch vectors are combined before classification. We intentionally avoid

architecture-specific pooling heads such as differentiable pooling or self-attention pooling (Ying et al., 2018; Lee et al., 2019), because the purpose is to isolate residual topology rather than to compare readout mechanisms. The classifier is trained with cross-entropy loss. Figure 1 summarizes the shared pipeline and the residual routes compared in the study.




Figure 1. Architecture overview of the compared residual neighborhoods on the branch-by-layer grid. The same message-passing backbone and graph-level readout are kept fixed while the residual routes are changed.

2.3 Residual neighborhoods

Plain. The plain baseline has no residual neighborhood:

$$\mathcal{N}(b, \ell) = \emptyset. \quad (3)$$

VerticalRes. Vertical residual reuse injects the previous layer on the same branch:

$$\mathcal{N}(b, \ell) = \{(b, \ell - 1)\}. \quad (4)$$

HorizontalRes. Horizontal residual reuse exchanges information across adjacent branches at the same depth:

$$\mathcal{N}(b, \ell) = \{(b - 1, \ell), (b + 1, \ell)\}, \quad (5)$$

with boundary branches using only valid neighbors. When $B = 1$, this horizontal neighborhood is empty.

MatrixRes. Matrix residual reuse combines vertical, horizontal, and diagonal neighborhoods:

$$\mathcal{N}(b, \ell) = \{(b, \ell - 1), (b - 1, \ell), (b + 1, \ell), (b - 1, \ell - 1), (b + 1, \ell - 1)\}. \quad (6)$$

This local stencil is intentionally not dense all-to-all mixing. It preserves an interpretable residual structure over the branch-by-layer grid. Invalid boundary neighbors are removed; therefore, when $B = 1$, MatrixRes keeps the same-branch vertical residual path after the first layer but has no branch-neighbor routes.

MatrixResGated. MatrixResGated uses the same neighborhood as MatrixRes, but the residual transform controls the injected signal:

$$\mathcal{R}(\mathbf{U}) = \begin{cases} \alpha \mathbf{U}, & \text{identity residual filtering with a gate,} \\ \alpha \text{softshrink}_{\lambda}(\mathbf{U}), & \text{sparse residual filtering with a gate.} \end{cases} \quad (7)$$

In the main benchmark, Plain, VerticalRes, HorizontalRes, and MatrixRes use identity residual filtering. MatrixResGated uses sparse residual filtering with $\lambda = 0.05$ and a learnable scalar gate initialized at 0.8. When $B = 1$, MatrixResGated still applies gated sparse filtering to the remaining matrix-residual paths. The implementation also supports top- k residual filtering and fixed gates, but these are not the default settings used for the main reported results.

165 2.4 Datasets and evaluation protocol

166 The main benchmark uses six datasets from TUDataset (Morris et al., 2020): PROTEINS,
 167 DD, ENZYMES, MUTAG, AIDS, and Mutagenicity. PROTEINS and DD are binary protein-
 168 structure graph classification benchmarks, ENZYMES is a six-class enzyme classification bench-
 169 mark, and MUTAG, AIDS, and Mutagenicity provide supplementary molecular-graph robust-
 170 ness checks. All datasets are loaded through PyTorch Geometric (Fey and Lenssen, 2019). All
 171 experiments use the raw benchmark features and graph structures without task-specific feature
 172 engineering, graph augmentation, or edge redesign.

173 All main results use graph-level five-fold evaluation. For fold k , the test accuracy is

$$\text{Acc}_k = \frac{1}{|\mathcal{T}_k|} \sum_{(\mathcal{G}_i, y_i) \in \mathcal{T}_k} \mathbf{1}[\hat{y}_i = y_i]. \quad (8)$$

174 Across folds, we report

$$\overline{\text{Acc}} = \frac{1}{5} \sum_{k=1}^5 \text{Acc}_k, \quad (9)$$

175 and the fold-level standard deviation.

176 2.5 Synthetic multi-class structural benchmark

177 To test residual topology under controlled structural class granularity, we add a synthetic bench-
 178 mark based on five graph families: Bernoulli random graphs recorded as ER, Barabási–Albert
 179 (BA) graphs, stochastic block models (SBM), Watts–Strogatz (WS) graphs, and a regular-
 180 degree control. For each family, we generate seven multi-class datasets with class counts
 181 $C \in \{2, 3, 4, 5, 6, 7, 8\}$. Each class contains 200 graphs, so each dataset contains $200C$ graphs.
 182 Graph sizes are sampled between 30 and 80 nodes, and every node receives an eight-dimensional
 183 feature vector derived from normalized degree and random feature channels. The features are
 184 intentionally generic and do not encode the class label directly.

185 The synthetic labels are controlled by family-specific structural parameters. The ER identi-
 186 fier in our experiment records corresponds strictly to Gilbert’s $G(n, p)$ model: each node pair is
 187 connected independently with Bernoulli probability p , producing sparse random structures at
 188 different densities (Gilbert, 1959). BA classes differ by the preferential-attachment count and
 189 provide long-tailed degree structures (Barabási and Albert, 1999). SBM classes vary the within-
 190 block and between-block probabilities to control community separation (Holland et al., 1983).
 191 WS classes differ by rewiring probability and provide small-world structures with local clus-
 192 tering and short paths (Watts and Strogatz, 1998). The regular-degree control uses randomly
 193 relabeled regular ring lattices with different fixed degrees. It provides degree-uniform structures
 194 but is not sampled from the uniform random-regular distribution, for which (Wormald, 1999)
 195 provides background. These families are used as stress tests rather than as replacements for real
 196 graph-classification benchmarks.

197 The synthetic benchmark uses the same five model families, four message-passing opera-
 198 tors, branch count $B = 3$, and five-fold protocol as the main benchmark, producing 35 datasets
 199 and 3,500 completed model runs. Because accuracy values are not directly comparable across
 200 different class counts, we report macro-F1 and chance-normalized accuracy in addition to raw
 201 accuracy. For a C -class dataset, normalized accuracy is

$$\text{NAcc} = \frac{\text{Acc} - 1/C}{1 - 1/C}. \quad (10)$$

202 This metric maps chance-level accuracy to 0 and perfect accuracy to 1, making the two- to
 203 eight-class scaling trend easier to interpret.

204 2.6 Training configuration

205 All main benchmark runs use the default branch count $B = 3$. The benchmark is repeated
 206 under four message-passing operators: GCNConv, GATConv, SAGEConv, and GINConv. The
 207 hidden dimension is 64 unless a sensitivity setting changes it. The default depth is $L = 4$

message-passing layers, except DD, where $L = 3$ is used for the main benchmark. All models use ReLU nonlinearities and the default PyTorch/PyTorch Geometric parameter initialization for the corresponding layers. Models are optimized with Adam (Kingma and Ba, 2014); the checkpoint with the lowest validation loss is retained for test evaluation. Early stopping monitors validation loss, and all runs use ReduceLROnPlateau with factor 0.5, patience 15, and minimum learning rate 10^{-5} . Gradients are clipped at norm 2.0. The branch-count ablation evaluates $B = 1, \dots, 8$ for HorizontalRes, MatrixRes, and MatrixResGated on PROTEINS and DD under GCNConv. The sensitivity scan evaluates MatrixRes and MatrixResGated on fold 0 at $B = 3$, varying learning rate, dropout, hidden dimension, sparsity strength, and gate initialization. Selected MatrixResGated candidates are then rerun across all five folds.

Table 1. Key experimental settings shared by the main benchmark.

Parameter	Value
Message-passing operators	GCNConv, GATConv, SAGEConv, GINConv
Default branch count	3
Default layers	4, except DD uses 3
Default hidden dimension	64
Evaluation	Five-fold graph classification
Synthetic stress test	ER, BA, SBM, WS, and regular-degree controls with $C = 2, \dots, 8$, 200 graphs per class
Branch-count ablation	$B = 1, \dots, 8$
Primary metrics	Mean best test accuracy; synthetic benchmark also reports macro-F1 and normalized accuracy
Sensitivity scope	Fold 0, $B = 3$
Plain/VerticalRes/HorizontalRes	200 epochs, patience 60, learning rate 0.005, weight decay 10^{-4} , dropout 0.5
MatrixRes/MatrixResGated	240 epochs, patience 80, learning rate 0.003, weight decay 5×10^{-5} , dropout 0.3
Dataset-specific settings	DD uses dropout 0.3; ENZYMES uses batch size 128
Learning-rate schedule	ReduceLROnPlateau, factor 0.5, patience 15, min lr 10^{-5}
Gradient clipping	Norm 2.0

2.7 Mechanism metrics

To interpret branch-count behavior, we track several mechanism indicators on the final branch-level graph embeddings. Branch diversity is the mean pairwise Euclidean distance between branch embeddings. Branch cosine similarity is the mean pairwise cosine similarity, and branch CKA is linear centered-kernel alignment computed on the same branch-embedding set. For $B = 1$, there is no branch pair; we report diversity as 0 and cosine and CKA as 1 under the single-branch self-similarity convention. Mean gradient norm is the representative parameter-gradient norm recorded during training. Residual-count totals summarize how many residual inputs are available across residual sites, active ratio measures the fraction of nonzero residual entries after filtering, and the gate statistic reports the learned scalar gate used by MatrixResGated. These metrics are used as correlational diagnostics rather than as causal tests.

3 RESULTS

All results are reported as five-fold mean best test accuracy $\pm$ fold-level standard deviation unless otherwise stated.

3.1 Overall benchmark performance

The main benchmark asks whether residual topology has a stable winner across datasets and message-passing operators. It does not. Table 2 summarizes the full benchmark: across the 24 dataset-operator combinations, MatrixRes is the most frequent winner, with 9 wins and the best mean rank (2.12). MatrixResGated and Plain each win 5 combinations, HorizontalRes wins

237 3, and VerticalRes wins 2. Thus, matrix-style reuse has the most favorable aggregate profile in
 238 this controlled benchmark, but every topology has a setting in which it is the strongest choice.

Table 2. Winner counts and mean ranks across the 24 dataset-operator combinations. Lower mean rank is better.

Model	Wins	Mean rank	Interpretation
Plain	5	3.04	Strong on selected small or easier settings
VerticalRes	2	3.38	Depth-only reuse helps selectively
HorizontalRes	3	3.54	Competitive when lateral exchange is sufficient
MatrixRes	9	2.12	Most frequent winner in this benchmark
MatrixResGated	5	2.92	Useful when sparse control helps residual traffic




Figure 2. Model-level winner counts across the 24 dataset-operator combinations.

239 Table 3 gives the cell-level evidence behind the win counts in Table 2. The matrix family
 240 is most prominent on MUTAG, AIDS, and Mutagenicity, where MatrixRes or MatrixResGated
 241 wins most operator settings. The protein-oriented datasets are more mixed: PROTEINS al-
 242 ternates between Plain and HorizontalRes, while DD alternates among HorizontalRes, MatrixRes,
 243 and VerticalRes. ENZYMES remains difficult and does not produce a stable residual-topology
 244 preference.

Table 3. Winning model for each dataset and message-passing operator. Values in parentheses are mean best test accuracy.

Dataset	GCNConv	GATConv	SAGEConv	GINConv
PROTEINS	Plain (0.7044)	HorizontalRes (0.7116)	Plain (0.6999)	HorizontalRes (0.7080)
DD	HorizontalRes (0.7181)	MatrixRes (0.7164)	MatrixRes (0.7198)	VerticalRes (0.7224)
ENZYMES	MatrixResGated (0.2933)	Plain (0.2717)	VerticalRes (0.2883)	Plain (0.3100)
MUTAG	MatrixRes (0.7609)	MatrixResGated (0.7555)	MatrixRes (0.7504)	MatrixResGated (0.8081)
AIDS	MatrixRes (0.8365)	MatrixRes (0.8875)	Plain (0.8850)	MatrixResGated (0.9280)
Mutagenicity	MatrixRes (0.7784)	MatrixRes (0.7909)	MatrixResGated (0.8031)	MatrixRes (0.8024)

245 The ENZYMES outcomes should be interpreted cautiously because the absolute accuracies
 246 are low for all compared models. Keeping ENZYMES in the full 24-combination summary
 247 preserves the preregistered benchmark scope, but it should not be read as strong evidence for
 248 a residual topology by itself. Excluding ENZYMES leaves MatrixRes with 9 wins across 20
 249 dataset-operator combinations and a mean rank of 2.05, so the aggregate MatrixRes preference

remains. The split is still context-dependent: MatrixRes is most prominent on the molecular datasets, whereas the protein-oriented datasets produce mixed winners.

3.2 GCNConv benchmark slice

To further remove variation caused by message-passing operators and observe the effect of residual topology under a fixed backbone, we report the GCNConv slice separately. Table 4 gives the full slice at the default branch count $B = 3$. Plain is strongest on PROTEINS, HorizontalRes is strongest on DD, MatrixResGated is strongest on ENZYMES, and MatrixRes is strongest on MUTAG, AIDS, and Mutagenicity.

Table 4. GCNConv benchmark results with default branch count $B = 3$.

Dataset	Plain	VerticalRes	HorizontalRes	MatrixRes	MatrixResGated
PROTEINS	0.7044 ± 0.0387	0.6993 ± 0.0340	0.6992 ± 0.0399	0.6973 ± 0.0324	0.6886 ± 0.0397
DD	0.7053 ± 0.0517	0.7143 ± 0.0260	0.7181 ± 0.0296	0.7127 ± 0.0365	0.7141 ± 0.0359
ENZYMES	0.2683 ± 0.0343	0.2550 ± 0.0455	0.2633 ± 0.0584	0.2750 ± 0.0577	0.2933 ± 0.0470
MUTAG	0.7556 ± 0.0439	0.7451 ± 0.0695	0.7343 ± 0.0433	0.7609 ± 0.0715	0.7556 ± 0.0721
AIDS	0.8215 ± 0.0046	0.8210 ± 0.0233	0.8130 ± 0.0297	0.8365 ± 0.0412	0.8150 ± 0.0231
Mutagenicity	0.7644 ± 0.0188	0.7667 ± 0.0212	0.7738 ± 0.0134	0.7784 ± 0.0185	0.7701 ± 0.0165




Figure 3. GCNConv benchmark at $B = 3$. Bars show five-fold mean best test accuracy with fold-level standard-deviation error bars.

The GCNConv slice supports the same conclusion as the full benchmark: residual topology matters, but the best topology depends on the dataset. MatrixRes is the strongest default choice on three of six GCNConv datasets, while simpler residual paths are still preferable on PROTEINS and DD. The ENZYMES winner is reported for completeness, but its low absolute accuracy limits the strength of that comparison.

3.3 Synthetic multi-class scaling

The synthetic benchmark asks a different question from the real-data benchmark: does residual topology behave differently when the structural class granularity is increased under controlled graph-generation rules? Table 5 summarizes the class-count trend after averaging over the five synthetic graph families and four message-passing operators. The two- and three-class settings do not favor matrix-style reuse: Plain is best at $C = 2$, and HorizontalRes is best at $C = 3$. From $C = 4$ onward, matrix-style reuse becomes more competitive. MatrixResGated is the strongest mean-accuracy model at $C = 4$, $C = 5$, $C = 6$, and $C = 8$, while VerticalRes is strongest at $C = 7$.

Table 5. Synthetic multi-class scaling averaged over five graph families and four operators. NAcc denotes chance-normalized accuracy.

Classes	Best model	Best Acc.	Plain Acc.	MatrixRes Acc.	MatrixResGated Acc.	Best NAcc
2	Plain	0.9585	0.9585	0.9481	0.9578	0.9170
3	HorizontalRes	0.8800	0.8774	0.8600	0.8772	0.8200
4	MatrixResGated	0.8111	0.8002	0.8083	0.8111	0.7481
5	MatrixResGated	0.7296	0.7233	0.7280	0.7296	0.6621
6	MatrixResGated	0.6878	0.6719	0.6827	0.6878	0.6254
7	VerticalRes	0.6524	0.6314	0.6429	0.6395	0.5945
8	MatrixResGated	0.6035	0.5865	0.6002	0.6035	0.5469



Figure 4. Synthetic multi-class scaling from two to eight classes. Each point averages five synthetic graph families and four message-passing operators.

Table 6 focuses on the higher-class regime $C = 4, \dots, 8$, where the task requires finer structural discrimination. MatrixResGated has the best mean accuracy and normalized accuracy in this regime, improving accuracy from 0.6827 for Plain to 0.6943. MatrixRes has the strongest macro-F1 among the compared models and is close to MatrixResGated in normalized accuracy. The class-scaling curve in Figure 4 also shows a stability pattern: from $C = 2$ to $C = 8$, MatrixRes loses 0.3479 accuracy points, compared with 0.3720 for Plain. Thus, the synthetic benchmark supports a narrow but useful interpretation: matrix-style reuse is not necessary for easy low-class settings, but it becomes more favorable as the structural labels become more fine-grained.

Table 6. Synthetic high-class regime averaged over $C = 4, \dots, 8$, five graph families, and four operators.

Model	Accuracy	Macro-F1	NAcc
Plain	0.6827	0.6703	0.6183
VerticalRes	0.6931	0.6834	0.6307
HorizontalRes	0.6846	0.6722	0.6210
MatrixRes	0.6924	0.6880	0.6300
MatrixResGated	0.6943	0.6876	0.6324

3.4 Branch-count ablation

Branch count is not merely a capacity parameter: it also changes residual-path density, so it requires a separate ablation. Table 7 reports the best and worst branch counts for each topology. On PROTEINS, the useful region is moderate: HorizontalRes peaks at $B = 4$, MatrixRes peaks at $B = 6$, and MatrixResGated peaks at $B = 3$. On DD, the useful region is smaller: HorizontalRes and MatrixRes peak at $B = 2$, while MatrixResGated is strongest already at $B = 1$. Thus, PROTEINS benefits from moderate branch budgets, whereas DD favors smaller budgets. Branch expansion must be matched to task structure.

Table 7. Best and worst branch-count settings in the ablation.

Dataset	Model	Best B	Best accuracy	Worst B	Worst accuracy
PROTEINS	HorizontalRes	4	0.7125 ± 0.0286	1	0.6801 ± 0.0518
PROTEINS	MatrixRes	6	0.7062 ± 0.0210	1	0.6909 ± 0.0586
PROTEINS	MatrixResGated	3	0.7098 ± 0.0219	4	0.6711 ± 0.0471
DD	HorizontalRes	2	0.7275 ± 0.0304	4	0.7071 ± 0.0302
DD	MatrixRes	2	0.7206 ± 0.0411	7	0.7071 ± 0.0390
DD	MatrixResGated	1	0.7283 ± 0.0244	7	0.7053 ± 0.0456




Figure 5. Branch-count ablation for HorizontalRes, MatrixRes, and MatrixResGated on PROTEINS and DD under GCNConv. The curves are not monotonic: accuracy improves only within an effective operating region before saturating or declining.

287 Table 7 and Figure 5 show that branch count should not be read as a simple capacity knob.
288 Increasing B changes the residual graph itself: it adds branches, residual paths, residual traffic,
289 and optimization interactions. PROTEINS tolerates moderate branch expansion, while DD is
290 more sensitive to redundant branch expansion.

291 3.5 Mechanism analysis: diversity, similarity, gradients, and residual traffic

292 To explain why branch-count changes affect performance rather than merely report accuracy,
293 we analyze branch diversity, branch cosine similarity, branch CKA similarity, representative
294 gradient norms, residual-count totals, active ratios, and gate values. Table 8 reports these
295 mechanism metrics at the best branch count for each ablation target. Diversity measures how
296 far apart branch representations are; branch cosine and branch CKA measure representational
297 redundancy or alignment; gradient norm captures optimization signal; residual count and active
298 ratio describe the amount of residual traffic injected into the model. These measurements are
299 diagnostic correlations rather than causal interventions.

Table 8. Mechanism metrics at the best branch count for each ablation target.

Dataset	Model	B	Acc.	Diversity	Cosine	CKA	Grad.	Residuals	Active	Gate
PROTEINS	HorizontalRes	4	0.7125	1.1450	0.8417	0.9797	0.0157	24.0	0.2519	0.7911
PROTEINS	MatrixRes	6	0.7062	3.7755	0.9173	0.9634	0.0182	88.0	0.4273	0.7757
PROTEINS	MatrixResGated	3	0.7098	1.3555	0.9790	0.9988	0.0316	37.0	0.3121	0.7833
DD	HorizontalRes	2	0.7275	0.1687	0.9969	1.0000	0.0732	6.0	0.3492	0.8375
DD	MatrixRes	2	0.7206	0.1834	0.9991	1.0000	0.0796	14.0	0.5384	0.8222
DD	MatrixResGated	1	0.7283	0.0000	1.0000	1.0000	0.0821	2.0	0.2352	0.8221

300 The PROTEINS ablation shows that moderate branch separation can be useful. Horizontal-
301 Res improves from $B = 1$ to $B = 4$ while branch diversity increases from 0.0000 to 1.1450 and
302 branch cosine decreases from 1.0000 to 0.8417. However, diversity alone does not guarantee bet-
303 ter accuracy: at $B = 8$, diversity rises further to 1.5441, but accuracy falls to 0.6881. MatrixRes

304 follows a similar pattern at a larger scale. It peaks at $B = 6$, where branch diversity reaches
305 3.7755, but adding still more branches does not improve accuracy.

306 DD is more conservative. HorizontalRes reaches its best result at $B = 2$, where branch
307 diversity is only 0.1687 and branch cosine remains 0.9969. MatrixRes also peaks at $B = 2$, with
308 similarly high branch cosine (0.9991). Larger branch counts increase diversity on DD, but the
309 performance does not improve. This pattern suggests that larger residual neighborhoods are
310 not useful on DD unless the added branches produce sufficiently useful functional separation.

311 The sparse/gated model clarifies the role of residual control. MatrixResGated has an active
312 ratio of 0.3121 at the best PROTEINS setting and 0.2352 at the best DD setting. These values
313 are consistent with sparse filtering suppressing residual entries rather than behaving as a dense
314 residual path. At the same time, MatrixResGated does not dominate MatrixRes overall, so
315 sparse control should be treated as a useful traffic-control mechanism rather than as a universal
316 improvement.




Figure 6. Mechanism summary linking branch-count accuracy to branch diversity, branch cosine similarity, branch CKA, and mean gradient norm.

317 3.6 MatrixRes and MatrixResGated sensitivity

318 The fold-0 sensitivity scan identifies candidate operating regions for MatrixRes and MatrixRes-
319 Gated. These scans are used for diagnosis rather than final claims. On PROTEINS, MatrixRes
320 does not improve over its fold-0 baseline in the tested sweep, while MatrixResGated has a fold-0
321 sparsity candidate: $\lambda = 0.02$ improves fold-0 accuracy from 0.6726 to 0.6951. Stronger sparsifica-
322 tion is harmful. On DD, both MatrixRes and MatrixResGated prefer lighter hidden dimensions
323 in the fold-0 scan: MatrixRes reaches 0.7595 with $d = 32$, and MatrixResGated reaches 0.7511
324 with $d = 32$. MatrixResGated also reaches 0.7511 with $lr = 0.001$ and $gate_init = 0.2$.

Table 9. Best fold-0 sensitivity settings for MatrixRes and MatrixResGated.

Dataset	Model	Baseline	Best setting	Best accuracy
DD	MatrixRes	0.7468	$dim = 32$	0.7595
DD	MatrixResGated	0.7300	$dim = 32$	0.7511
PROTEINS	MatrixRes	0.6637	baseline	0.6637
PROTEINS	MatrixResGated	0.6726	$\lambda = 0.02$	0.6951

MatrixResGated sensitivity slices




Figure 7. Fold-0 MatrixResGated sensitivity scan at $B = 3$. The scan identifies candidate operating regions but is not treated as final hyperparameter evidence.

3.7 Five-fold check of tuned candidates

Fold-0 scans are used only for candidate discovery and cannot establish stable performance by themselves. To obtain more reliable evidence, selected MatrixResGated settings were rerun across all five folds. Table 10 reports these follow-up checks. On DD, the MatrixResGated candidate with hidden dimension 32 reduces parameters from 42,371 to 15,043 while remaining competitive, indicating better parameter efficiency. The lower learning rate is close to the default MatrixResGated benchmark, and the lower gate initialization is weaker. On PROTEINS, the candidate with sparsity strength $\lambda = 0.02$ does not consistently preserve its fold-0 advantage across all folds, so the sparsity gain should not be overstated.

Table 10. Five-fold checks of selected MatrixResGated candidates.

Candidate	Dataset	Accuracy	Mean loss	Parameters
Hidden dimension 32	DD	0.7215 ± 0.0330	0.5717	15,043
Gate initialization 0.2	DD	0.7131 ± 0.0213	0.5820	42,371
Learning rate 0.001	DD	0.7181 ± 0.0336	0.5737	42,371
Sparsity strength $\lambda = 0.02$	PROTEINS	0.6909 ± 0.0447	0.6104	38,339

The tuned-candidate check prevents over-claiming from fold-0 sensitivity. It supports a smaller, parameter-efficient MatrixResGated configuration on DD, but it does not support a strong full-fold sparsity claim on PROTEINS.

4 DISCUSSION

This section answers the research questions in light of the six-dataset, four-operator benchmark, the branch-count ablation, the tuned-candidate checks, the mechanism evidence, and the synthetic multi-class scaling benchmark.

341 4.1 RQ1: Does residual topology affect performance under a shared graph-classification proto-
342 col?

343 The main benchmark argues against a universal residual topology. Across 24 dataset-operator
344 combinations, MatrixRes is the most frequent winner and has the best mean rank, but all
345 five model families win at least two combinations. Under GCNConv alone, Plain is strongest
346 on PROTEINS, HorizontalRes is strongest on DD, MatrixResGated is strongest on ENZYMES,
347 and MatrixRes is strongest on MUTAG, AIDS, and Mutagenicity. These outcomes indicate that
348 the useful residual neighborhood depends on dataset-specific tolerance for branch specialization,
349 residual traffic, and optimization complexity.

350 This conclusion is consistent with the broader GNN literature: changing the propagation
351 operator, normalization, or pooling design often shifts the preferred architecture (Xu et al., 2019;
352 Hamilton et al., 2017; Veličković et al., 2018; Ying et al., 2018). The present study adds that
353 residual topology itself is another source of this shift. Matrix-style reuse can be read as a local
354 residual stencil over the branch-by-layer grid, analogous in spirit to structured neighborhood
355 mixing (Abu-El-Haija et al., 2019), but applied to hidden-state reuse rather than to graph
356 adjacency. Its advantage is therefore architectural organization, not a new message-passing
357 primitive.

358 Answer: MatrixRes is the most frequent winner in the current full benchmark, but residual
359 topology is still dataset- and operator-dependent. A graph-classification benchmark should
360 compare vertical, horizontal, matrix-style, and sparse/gated reuse rather than assuming that
361 one residual direction is universally best.

362 4.2 RQ2: How do branch count, residual traffic, and mechanisms explain performance changes?

363 The branch-count ablation shows that B is not only a capacity parameter. Increasing B adds
364 branches, residual paths, and optimization interactions. In the GCNConv ablation, PROTEINS
365 has a moderate useful region: HorizontalRes peaks at $B = 4$, MatrixRes at $B = 6$, and MatrixRes-
366 Gated at $B = 3$. DD prefers smaller budgets, with peaks at $B = 1$ or $B = 2$. The mechanism
367 summaries are consistent with this split. On PROTEINS, MatrixRes with $B = 6$ produces the
368 largest branch diversity, whereas HorizontalRes keeps lower branch cosine and a smaller resid-
369 ual footprint. On DD, all best branch-count settings have nearly saturated branch cosine and
370 CKA, so extra residual traffic is less likely to create new useful functions. Branch diversity may
371 increase with B , but if branch similarity remains high or gradients weaken, the extra branches
372 do not translate into better classification.

373 This rise-then-fall behavior is also a practical warning for residual GNN design. Residual
374 connections are often introduced to make deeper or wider models easier to optimize (He et al.,
375 2016; Li et al., 2019, 2021), but a multi-branch residual model changes both capacity and
376 routing density. A larger B increases the number of paths through which hidden states can
377 be reinjected, so the useful regime is governed by the balance between functional diversity
378 and residual interference. The branch-count results suggest that this balance differs sharply
379 between PROTEINS and DD, even though both are commonly treated as protein-structure
380 graph benchmarks.

381 Answer: In the PROTEINS/DD GCNConv ablation, branch count helps only within a
382 limited operating region. Larger branch budgets are useful when they create functional diversity,
383 and less useful when residual complexity grows faster than branch separation or optimization
384 quality.

385 4.3 RQ3: When is matrix-style residual reuse more effective on real data and controlled synthetic
386 multi-class structures?

387 MatrixResGated tests whether matrix residual traffic should be filtered before injection. The
388 main benchmark shows that sparse/gated reuse is not uniformly better than dense matrix reuse:
389 MatrixRes wins more dataset-operator combinations overall. However, MatrixResGated is still
390 useful because it wins several combinations and provides a controlled way to reduce weak residual
391 coefficients. The ENZYMES GCNConv win is reported but should be interpreted cautiously
392 because all compared models have low absolute accuracy on that dataset. The mechanism
393 table supports the residual-control interpretation: the gated variant generally keeps fewer active
394 residual pathways than dense MatrixRes, while the learned gate remains in a moderate range

rather than collapsing to zero. This distinction is important. The value of MatrixResGated is not that sparsity always improves accuracy, but that sparse residual control can be beneficial when matrix-style reuse would otherwise inject too much low-value residual traffic.

This places sparse/gated residual reuse between two familiar design ideas. Like residual gated graph networks, it learns how strongly historical information should be injected (Bresson and Laurent, 2017); like sparsification and regularization methods, it reduces the influence of weak or noisy pathways (Rong et al., 2019; Cangea et al., 2018). The present implementation is deliberately simple: it uses a scalar gate and soft-shrink filtering rather than a large routing network. This keeps the ablation interpretable and prevents the gated variant from becoming a separate high-capacity model.

MatrixResGated defines a tunable family rather than a plug-and-play winner. The fold-0 sensitivity scan suggests candidate operating regions: PROTEINS has a mild-sparsity candidate, while DD has conservative candidates based on smaller hidden dimension and lower initial gate. The five-fold candidate reruns refine this interpretation. On DD, the smaller hidden dimension remains beneficial and parameter-efficient. On PROTEINS, the mild-sparsity candidate does not preserve its fold-0 advantage across all folds.

Real-data evidence. Sparse/gated matrix reuse is a control mechanism for residual traffic. It is complementary to dense MatrixRes rather than a strict replacement, and candidate settings must be validated across folds before they are treated as performance claims. The current evidence supports a smaller controlled matrix model on DD and leaves PROTEINS sparsification as an unstable candidate rather than a confirmed gain.

The synthetic multi-class benchmark separates class granularity from real-dataset idiosyncrasies. In the two- and three-class settings, simpler residual choices remain competitive: Plain is best at $C = 2$, and HorizontalRes is best at $C = 3$. This result is important because it prevents an overbroad claim that matrix-style reuse is always needed. When the class count increases, the pattern changes. In the $C = 4, \dots, 8$ regime, MatrixResGated has the best average accuracy and normalized accuracy, while MatrixRes has the best macro-F1. MatrixRes also has the smallest accuracy drop from $C = 2$ to $C = 8$, suggesting that dense local matrix reuse is relatively stable as the label space becomes finer.

The distribution-level synthetic results refine this conclusion. MatrixRes is strongest in the higher-class BA, ER, and regular-degree control settings, where class labels are tied to degree growth, edge probability, or degree-uniform structural changes. The advantage is weaker on SBM and absent on WS, where community structure and small-world rewiring create different inductive pressures. Therefore, the synthetic benchmark supports a conditional claim: matrix-style residual reuse is most useful when the task requires separating several related structural regimes, but it does not solve every graph distribution.

Answer: The real-data and synthetic scaling experiments jointly show that MatrixRes and MatrixResGated are conditionally effective matrix-residual families rather than universal replacements. Sparse/gated reuse is useful for controlling excessive residual traffic, while matrix-style reuse becomes more favorable under finer structural discrimination, especially from four to eight classes. The effect is strongest on several controlled degree- and density-driven graph families, but it remains distribution-dependent.

4.4 Limitations

Several limitations bound the generality of these findings. First, although the main benchmark includes four message-passing operators, it still uses a fixed graph-level pooling and classification head; alternative readout designs may interact with residual topology. Second, the compared operators are canonical message-passing backbones rather than recent graph transformer or local-global hybrid architectures (Dwivedi and Bresson, 2020). This is a deliberate control choice, but it means the paper should not be read as a state-of-the-art model comparison. Third, the sensitivity scan is fold-0 only; although selected candidates were rerun across five folds, that follow-up is limited to MatrixResGated on PROTEINS and DD. Fourth, ENZYMES has low absolute accuracy and high variability, so its result should be interpreted cautiously. Finally, the mechanism analysis is correlational. Branch diversity, cosine similarity, CKA, and gradient norms support a coherent explanation, but they do not prove causality.

Another limitation is that the benchmark intentionally avoids task-specific chemical or biological feature engineering. This makes the residual-topology comparison cleaner, but it also means that the absolute accuracies should not be interpreted as optimized state-of-the-art molecular prediction results. The synthetic benchmark has a related limitation: its labels are generated from controlled structural parameters, so it is useful for stress testing residual topology but cannot substitute for broader real multi-class graph benchmarks. The paper’s contribution is methodological: it identifies when different residual pathways are helpful under matched conditions.

4.5 Future directions

Future work should extend the comparison to larger graph-classification datasets, richer readout mechanisms, recent local-global backbones, and adaptive residual-neighborhood selection. A natural next step is to replace manually chosen vertical, horizontal, or matrix stencils with a learned local residual policy. The tuned-candidate validation should also be broadened beyond the current MatrixResGated checks to determine whether dataset-specific residual control can be selected reliably.

5 CONCLUSIONS

This paper studies residual connectivity for graph classification as a structured branch-by-layer design problem. Across six datasets and four message-passing operators, residual topology matters, but its value depends on the dataset, operator, and branch budget. MatrixRes is the most frequent winner and has the best mean rank in the full benchmark, but the per-dataset GCNConv winners remain mixed: Plain is strongest on PROTEINS, HorizontalRes on DD, MatrixResGated on ENZYMES, and MatrixRes on MUTAG, AIDS, and Mutagenicity. The ENZYMES comparison has low absolute accuracy and should be interpreted cautiously. The branch-count ablations further show that adding branches is useful only within a limited operating region.

The mechanism analysis provides a plausible explanation for why branch expansion is not automatically beneficial. Additional branches can increase representational diversity, but they also increase residual traffic and may preserve redundant branch similarity or weaken optimization. Matrix-style reuse is best viewed as a controlled residual family whose effectiveness depends on the branch budget and residual filtering strategy.

The synthetic multi-class benchmark strengthens this interpretation by varying structural class granularity under controlled graph-generation rules. Matrix-style reuse is not required for the easiest two- and three-class settings, but it becomes more competitive from four to eight classes. In that higher-class regime, MatrixResGated has the best average accuracy and normalized accuracy, while MatrixRes has the smallest two- to eight-class accuracy drop. This supports the view that matrix residual neighborhoods are most useful when graph-level classification requires finer structural discrimination.

The practical implication is that residual topology, branch count, operator choice, and residual control should be evaluated jointly. Before treating a sensitivity result as a performance claim, candidate settings should be rerun across folds. In the current experiments, this validation supports a smaller and more parameter-efficient MatrixResGated model on DD, while the PROTEINS sparsity candidate remains less stable.

6 ACKNOWLEDGMENTS

The authors thank the open-source graph learning community and the maintainers of the benchmark resources used in this study. They also thank the reviewers and editors for their time and feedback.

A APPENDIX

A.1 Result sources

The manuscript uses the following summary files as the source of truth. The main benchmark summary contains 120 model-level rows: six datasets, four operators, five model families, and

500 five-fold aggregation for each combination.

- 501 • records/LATEST/summaries/benchmark_summary.csv
- 502 • records/LATEST/summaries/branch_ablation_summary.csv
- 503 • records/LATEST/summaries/parameter_sensitivity_summary.csv
- 504 • records/LATEST/summaries/tuned_candidate_summary.csv
- 505 • records/LATEST/summaries/mechanism_compact_summary.csv
- records/SYN_MULTI_FULL/summaries/benchmark_summary.csv

506 A.2 Supplementary diagnostics

507 Additional fold-level rows, CKA summaries, residual traffic summaries, and full sensitivity tables
508 are retained in the repository records directory for supplementary conversion.

509 A.3 Evidence map

510 Table 11 summarizes where each main empirical claim is supported in the manuscript. This
511 compact table is included to keep the main text readable while still making the result-to-claim
512 link explicit.

Table 11. Compact map from empirical claims to manuscript evidence.

Research question	Claim	Primary evidence
RQ1	MatrixRes is the most frequent winner in the full benchmark but not a universal winner; GCNConv results remain dataset-dependent under a fixed operator.	Table 2, Table 3, Figure 2, Table 4, and Figure 3.
RQ2	Branch count produces a rise-then-fall pattern, and useful expansion is associated with separation without excessive residual traffic.	Table 7, Figure 5, Table 8, and Figure 6.
RQ3	Fold-0 sensitivity should not be treated as final evidence; matrix-style reuse becomes more favorable in higher-class synthetic structural discrimination.	Table 9, Figure 7, Table 10, Table 5, Figure 4, Table 6, and Table 12.

513 A.4 Synthetic distribution-level summary

514 Table 12 reports the higher-class synthetic results by graph family. The table is placed in
515 the appendix because the main text focuses on the class-count scaling trend. It shows that
516 MatrixRes is strongest on BA, ER, and the regular-degree control, while SBM and WS remain
517 less favorable to matrix-style reuse.

Table 12. Synthetic high-class accuracy by graph family, averaged over $C = 4, \dots, 8$ and four operators.

Family	Plain	VerticalRes	HorizontalRes	MatrixRes	MatrixResGated
BA	0.7874	0.8177	0.7855	0.8271	0.8257
ER	0.8777	0.8756	0.8728	0.8807	0.8796
Regular	0.7505	0.7923	0.7739	0.8077	0.8057
SBM	0.6268	0.6286	0.6218	0.6264	0.6176
WS	0.3709	0.3511	0.3693	0.3202	0.3430

518 REFERENCES

- 519 Sami Abu-El-Haija, Bryan Perozzi, Aalok Kapoor, Navid Alipourfard, Kristina Lerman, Greg
520 Steeg, and Aram Galstyan. Mixhop: Higher-order graph convolutional architectures via spar-
521 sified neighborhood mixing. In International Conference on Machine Learning, pages 33–42,
522 2019.
- 523 Uri Alon and Eran Yahav. On the bottleneck of graph neural networks. In International
524 Conference on Learning Representations (ICLR), 2021.
- 525 Albert-László Barabási and Réka Albert. Emergence of scaling in random networks. *Science*,
526 286(5439):509–512, 1999. doi: 10.1126/science.286.5439.509. URL [https://doi.org/10.1126/](https://doi.org/10.1126/science.286.5439.509)
527 [science.286.5439.509](https://doi.org/10.1126/science.286.5439.509).
- 528 Filippo Maria Bianchi, Franco Grattarola, and Cesare Alippi. Graph neural networks with con-
529 volutional arma filters. In IEEE Transactions on Pattern Analysis and Machine Intelligence,
530 2020.
- 531 Xavier Bresson and Thomas Laurent. Residual gated graph convnets. arXiv preprint
532 arXiv:1711.07553, 2017.
- 533 Cătălina Cangea, Petar Veličković, Nikola Jovanović, Thomas Kipf, and Pietro Liò. Towards
534 sparse hierarchical graph classifiers. In Pacific-Asia Conference on Knowledge Discovery and
535 Data Mining, pages 165–176, 2018.
- 536 David K Duvenaud, Dougal Maclaurin, Jorge Iparraguirre, Rafael Bombarell, Timothy Hirzel,
537 Alan Aspuru-Guzik, and Ryan P Adams. Convolutional networks on graphs for learning
538 molecular fingerprints. In Advances in neural information processing systems, pages 2224–
539 2232, 2015.
- 540 Vijay Prakash Dwivedi and Xavier Bresson. A generalization of transformer networks to graphs.
541 arXiv preprint arXiv:2012.09699, 2020.
- 542 Matthias Fey and Jan Eric Lenssen. Fast graph representation learning with pytorch geometric.
543 arXiv preprint arXiv:1903.02428, 2019.
- 544 Edgar N. Gilbert. Random graphs. *The Annals of Mathematical Statistics*, 30(4):1141–1144,
545 1959. doi: 10.1214/aoms/1177706098. URL <https://doi.org/10.1214/aoms/1177706098>.
- 546 Justin Gilmer, Samuel S Schoenholz, Patrick F Riley, Oriol Vinyals, and George E Dahl. Neural
547 message passing for quantum chemistry. In International Conference on Machine Learning,
548 pages 1263–1272, 2017.
- 549 Marco Gori, Gabriele Monfardini, and Franco Scarselli. A new model for learning in graph
550 domains. *Proceedings of IJCNN*, 2005.
- 551 William L Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large
552 graphs. In Advances in neural information processing systems, pages 1024–1034, 2017.
- 553 Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image
554 recognition. In Proceedings of the IEEE conference on computer vision and pattern recogni-
555 tion, pages 770–778, 2016.
- 556 Paul W. Holland, Kathryn Blackmond Laskey, and Samuel Leinhardt. Stochastic blockmodels:
557 First steps. *Social Networks*, 5(2):109–137, 1983. doi: 10.1016/0378-8733(83)90021-7. URL
558 [https://doi.org/10.1016/0378-8733\(83\)90021-7](https://doi.org/10.1016/0378-8733(83)90021-7).
- 559 Nicolas Keriven, Gabriel Peyré, and Samuel Vaiter. A unified benchmark for the diagnosis and
560 treatment of graph neural networks. arXiv preprint arXiv:2012.01450, 2020.
- 561 Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. arXiv preprint
562 arXiv:1412.6980, 2014.

563 Thomas N. Kipf and Max Welling. Semi-supervised classification with graph convolutional
564 networks. In International Conference on Learning Representations, 2017. URL [https://](https://arxiv.org/abs/1609.02907)
565 arxiv.org/abs/1609.02907.

566 Junhyun Lee, Inyeop Lee, and Jaewoo Kang. Self-attention graph pooling. arXiv preprint
567 arXiv:1904.08082, 2019.

568 Guohao Li, Matthias Müller, Ali Thabet, and Bernard Ghanem. DeepGCNs: Can GCNs go as
569 deep as CNNs? In IEEE/CVF International Conference on Computer Vision, 2019. URL
570 <https://arxiv.org/abs/1904.03751>.

571 Guohao Li, Elias Müller, Abbas Thabet, and Bernard Ghanem. Training graph neural networks
572 with 1000 layers. In International Conference on Machine Learning, pages 6403–6413, 2021.

573 Qimai Li, Zhichao Han, and Xiao-Ming Wu. Deeper insights into graph convolutional networks
574 for semi-supervised learning. In Proceedings of the AAAI Conference on Artificial Intelligence,
575 2018. URL <https://ojs.aaai.org/index.php/AAAI/article/view/11604>.

576 Christopher Morris, Nils M. Kriege, Franka Bause, Kristian Kersting, Petra Mutzel, and Marion
577 Neumann. TUDataset: A collection of benchmark datasets for learning with graphs. In ICML
578 Workshop on Graph Representation Learning and Beyond, 2020. URL [https://chrsmrrs.](https://chrsmrrs.github.io/datasets/)
579 [github.io/datasets/](https://chrsmrrs.github.io/datasets/).

580 Kenta Oono and Taiji Suzuki. Simple and deep graph convolutional networks. arXiv preprint
581 arXiv:2006.13604, 2020.

582 Yu Rong, Wenbing Huang, Tingyang Xu, and Junzhou Huang. Dropedge: Towards deep graph
583 convolutional networks on node classification. arXiv preprint arXiv:1907.10903, 2019.

584 Jake Topping, Francesco Di Giovanni, Benjamin Chamberlain, Michael Bronstein, Douwe Ven-
585 drig, and Pietro Lio. Understanding over-squashing and bottlenecks on graphs via curvature.
586 arXiv preprint arXiv:2111.14522, 2021.

587 Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Liò, and Yoshua
588 Bengio. Graph attention networks. In International Conference on Learning Representations
589 (ICLR), 2018.

590 Duncan J. Watts and Steven H. Strogatz. Collective dynamics of ‘small-world’ networks. *Nature*,
591 393:440–442, 1998. doi: 10.1038/30918. URL <https://doi.org/10.1038/30918>.

592 Nicholas C. Wormald. Models of random regular graphs. In J. D. Lamb and D. A. Preece, editors,
593 *Surveys in Combinatorics*, 1999, volume 267 of London Mathematical Society Lecture Note
594 Series, pages 239–298. Cambridge University Press, 1999. doi: 10.1017/CBO9780511721335.
595 010. URL <https://doi.org/10.1017/CBO9780511721335.010>.

596 Zonghan Wu, Shirui Pan, Fengwen Chen, Guoding Long, Cheng Zhang, and Sheng Yu Zhang.
597 A comprehensive survey on graph neural networks. *IEEE Transactions on Neural Networks*
598 *and Learning Systems*, 32(1):4–24, 2021.

599 Keyulu Xu, Chengtao Li, Yonglong Tian, Xiao Du, Jure Leskovec, and Stefanie Jegelka. Repre-
600 sentation learning on graphs with jumping knowledge networks. In International Conference
601 on Machine Learning, pages 5449–5458, 2018.

602 Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural
603 networks? In International Conference on Learning Representations, 2019. URL [https://](https://openreview.net/forum?id=ryGs6iA5Km)
604 openreview.net/forum?id=ryGs6iA5Km.

605 Rex Ying, Jiaxuan You, Christopher Morris, Xiang Ren, William L Hamilton, and Jure Leskovec.
606 Hierarchical graph representation learning with differentiable pooling. In *Advances in neural*
607 *information processing systems*, pages 4800–4810, 2018.